

DECISION NARRATIVE · CONFIDENTIAL

From Market Data to an AI Voice Front Desk

How we found the opportunity, pressure-tested it, and built the plan — the full evidence trail.

Prepared for: **My Partner**

Prepared by: Jaydeep (Founder) · June 2026

Companion files: SaaS Gap Research Data Bible (21 tabs) · Build Bible

Status: research complete · plan locked · ready for Phase 0

Executive Summary

We did not start with an idea and look for data to justify it. We started with a question — *"where is a real, winnable software gap we can build?"* — and ran a rigorous, multi-stage research process to answer it. This document is the evidence trail of how we got from raw market data to a concrete, fundable plan.

6

DATASETS ANALYSED

560

KEYWORDS SCORED

164

COMPETITORS MAPPED

What we concluded

The data pointed away from crowded, un-winnable categories (CRM, project management — SEO walls of 76–93) and toward a specific, high-value, under-served opportunity: **an AI voice "front desk" for appointment-driven local businesses** — software that answers the phone, qualifies callers, and books appointments, 24/7. It sits on the single most open keyword in our entire dataset (*"cloud based business phone system"* — SEO difficulty 9, \$76 cost-per-click) and on top of the most-hated, fastest-declining incumbents.

Why we can win it

We are not starting cold. We own **7 physiotherapy clinics, a psychology clinic, and work with construction/general-contractor clients** — ready-made design partners and first customers. We have a marketing agency (PPC Guru) that generates live inbound leads to test on, deep domain knowledge, and a lean, AI-agent-leveraged operating model that lets a tiny team move fast.

The ask

This is a capital-light start: roughly **\$5,000–\$15,000** takes us to a first paying external customer. We bootstrap on our own businesses' cash to product-market fit, then raise from strength. The next step is **Phase 0** — foundation, compliance-by-design, and the first pilot on our own agency's leads.

The Origin & The Question

The journey began with a deliberately open question and a pile of real research data — not a pet idea.

We assembled what we already had: a marketing agency running paid ads, an operations prototype (PPC Guru), a portfolio of service businesses, and a desire to build a SaaS product with a genuine, defensible gap. Rather than guess, we committed to letting the market data decide.

The founding question: "Given everything happening in SaaS and AI right now, where is a gap we can realistically build into — one with real demand, weak incumbents, and a path we can actually win?"

Our starting assets (the unfair advantages we already held)

- **Distribution:** 7 physiotherapy clinics + a psychology clinic + construction/GC clients + a network of realtors and mortgage brokers.
- **A live testing ground:** PPC Guru generates real inbound leads daily via paid ads.
- **Domain expertise:** 10+ years in performance marketing, Google/Meta ads, and agency operations.
- **An operating model:** a lean team (one developer, Shrikaanth) leveraged by AI agents acting as executors — fast, low-overhead build capacity.

These assets shaped the whole analysis: we weren't looking for *any* opportunity, we were looking for the one where *our* advantages compound.

Our Method — How We Did the Research

The credibility of the conclusion rests on the rigor of the method. Here is exactly what we did, in order.

Stage	What we did	Output
1. Demand data	Pulled 4 years of Google Keyword Planner search data across USA, Canada, and Worldwide — volume, trend, seasonality, cost-per-click, competition.	3 demand datasets
2. Supply data	Built a competitive teardown of 164 SaaS tools using two independent AI engines (for cross-validation) — valuation, GTM, architecture, and customer complaints.	2 supply datasets
3. The Data Bible	Computed everything into a 21-tab workbook: rankings, 4-year year-over-year, seasonality, category rollups, competitor complaint-coding, and a demand×supply gap score.	Data Bible v1–v2
4. The correction	Caught that ~84% of the search volume was branded (people typing known product names), not true demand. Re-ran research on non-branded, problem-based keywords.	418 → 560 keyword set
5. Winnability	Added SEO difficulty (Ubersuggest) to every keyword — so we could see not just demand, but what we could realistically <i>win</i> .	Opportunity Score

Why two AI engines? We deliberately ran the competitor teardown twice and only fully trusted findings where both engines independently agreed — a built-in cross-validation that filters out hallucination and bias. The most corroborated finding (across both) became our highest-confidence signal.

Every number in this narrative traces back to the committed **Data Bible** workbook — it is auditable, not anecdotal.

The Demand Side

25.1M

SEARCHES/MO WORLDWIDE

~84%

WAS BRANDED, NOT REAL DEMAND

SEO 8-35

THE WINNABLE GREY SPACE

1. It's a power-law market

A handful of terms dominate — the top 12 US keywords account for ~68% of volume. The giant head terms (*crm, notion, booking, leads*) are **SEO walls of 76–93**: impossible for a newcomer to win.

2. The data was hiding the real demand

Our most important methodological catch: most of the "demand" was **brand-navigation** — people typing the name of a product that already won. That tells you who's famous, not where the unmet need is. Stripping the brands out revealed the true picture.

3. The grey space appeared

Once we layered SEO difficulty on top, a clear pattern emerged: **real category demand at low difficulty** — winnable terms in the SEO 8–35 range with \$60–130 cost-per-click (high commercial intent, almost nobody ranking). The standout: "*cloud based business phone system*" — **6,600 searches, SEO difficulty 9, \$76 CPC**. That combination did not exist anywhere else in 560 keywords.

"Don't chase the famous category words — that's the founder's classic mistake. Win the winnable long-tail where demand meets low difficulty."

The Supply Side — Weak, Legacy Incumbents

A market gap needs two things: unmet demand *and* incumbents you can beat. The competitor data delivered both.

~70%

INCUMBENTS ARE LEGACY

#1

COMPLAINT: BAD DATA / PRICING

Declining

ZENDESK -24%, FRESHDESK -45%

The structural weakness

Roughly **70% of the 164 tools are "legacy" architectures with AI bolted on afterward** — not built AI-native. Their most repeated customer complaints clustered into four themes: **pricing that punishes growth, bad/unverified data, tool sprawl (narrow tools), and complexity.**

The unowned middle

We also found a structural hole: traditional CRMs stop at "deal won"; project tools are internal-only and don't understand customers. The **operational layer where customer work actually happens** — answering, scheduling, following up — is under-served by modern, AI-native software.

Crossed with the demand findings, one territory lit up brightest: **the live, voice-first customer-intake layer** — exactly where the highest CPCs, the most open keywords, and the most-declining incumbents (Zendesk, Freshdesk, dumb VoIP, human answering services) all converged.

The Pivot to Voice

We deliberately stepped back and asked the question a top investor would: ignoring our background entirely, what is the single best, AI-native, lowest-effort / highest-result thing to build right now?

The three signals the data screamed

1. **Customer support / intake** carried the highest CPCs in the entire dataset (\$80–126) and the most-declining incumbents.
2. **The phone layer** sat on the single most open keyword (SEO difficulty 9) with strong commercial value.
3. **"Why now"**: real-time voice AI only crossed the quality threshold in **2025** — natural turn-taking, interruption handling, and latency finally became good enough to replace a human on the phone.

The insight: local service businesses lose real money every day on **missed calls** — a missed call is a lost \$200–\$15,000 customer. An AI that answers 24/7, qualifies, and books has an ROI you can prove in week one. You're not selling "software" — you're selling **recovered revenue**.

This is also the rare AI opportunity where **lowest effort meets maximum result**: the hard real-time engineering can be rented (it's commoditizing), so a small team ships a real product in weeks, not years — while the defensible moat (vertical depth, data, integrations) is built on top.

What We're Building

An AI voice front desk for appointment-driven local businesses. One engine, two modes.

Model 1 — Inbound Receptionist

Answers calls *to* the business 24/7, answers questions from the business's own knowledge, qualifies the caller, and books the appointment directly into their calendar — with an SMS confirmation and a human handoff when needed.

Model 2 — Speed-to-Lead Caller

When a lead submits a form on a paid ad, the AI calls them back **within 60 seconds**, qualifies them, and books the meeting. This matters enormously: contacting a fresh lead within ~1 minute versus 30 minutes lifts conversion by roughly **390%** (the "5-minute rule"). No human sales team can dial that fast, 24/7.

One engine: both modes share the same brain, knowledge base, and booking logic. We build it once and point it at inbound calls, outbound lead-callbacks, or both.

What the AI owns vs. the human

The AI owns the entire top of the funnel — instant response, qualification, FAQs, booking, follow-up, and nurture (80% of a rep's time, and the part humans do worst). Humans close the high-value, consultative deals. Over time, as the system learns from our best calls, it closes more itself.

Market & Opportunity

Layer	Size (US + Canada, estimate)	Meaning
Beachhead — physiotherapy	~40,000 clinics · ~\$190–220M	Where we start (we own 7)
Adjacent healthcare	~450,000+ locations · ~\$2B+	Dental, chiro, vet, med-spa...
All appointment/phone-driven local business	2–3M+ businesses · \$10–15B+ TAM	The platform (3–4× globally)

Can it be a billion-dollar company?

Physiotherapy alone is a great ~\$20–40M revenue niche — not a unicorn. But the **multi-vertical local AI front desk is plausibly a billion-dollar company**. The comparables already exist: **Podium** (~\$3B valuation), **Weave** (healthcare comms, public), **ServiceTitan** (home-services vertical SaaS, ~\$9B). Local-SMB vertical software produces multi-billion outcomes — and AI is the wedge to take share from the legacy versions.

The math: \$100M ARR ≈ ~20,000 locations — a tiny fraction of a multi-million-business market. The path to \$1B is "nail physio → replicate the vertical playbook." It is an execution problem, not a market-size problem.

All figures are planning estimates derived from the research; to be validated, not treated as audited.

Why Us — The Unfair Advantages

Most founders chasing this would start cold. We start with a stacked deck.

Advantage	Why it matters
7 physio clinics + psychology clinic	Instant design partners + first customers + real call data, with zero sales friction.
Construction / GC clients	A ready-made <i>second</i> vertical (home services) to prove the model is repeatable.
PPC Guru marketing agency	A live source of inbound leads to pilot the speed-to-lead model on our own money.
Realtor / mortgage network	~10,000 realtors reachable — owner-direct distribution for the next verticals.
The research itself	A 21-tab evidence base most founders never build — we know exactly where to fish.
Lean AI-agent operating model	One developer + AI executor agents = fast, low-overhead build capacity.

These advantages compound: our clinics generate the training data that powers the product, which makes onboarding the next clinic faster, which makes the moat deeper. The flywheel starts spinning on day one because we own the first customers.

The Moat — Why This Gets Harder to Beat Over Time

The honest question every investor asks: "Can't a big AI company or a funded startup just crush this?" The answer is no — and here's the architecture of why.

Foundation-model companies (OpenAI, Google, Anthropic) sell the *model*; they don't build vertical, compliance-bound, integration-heavy applications. As models get cheaper, **our costs fall while our moat stays**. The real defensibility is built from things they will never do per-vertical:

- **Vertical depth + integrations** — booking into the clinic's practice-management system, insurance and intake workflows, recalls. Slow, unglamorous, and exactly what generic players skip.
- **Proprietary call data** → an **"Industry Brain"** — every call (anonymized) trains a per-vertical model. Each new clinic makes the product smarter for the next: a **data network effect** nobody starting fresh can match.
- **Accent adaptation** — a model fine-tuned on our own accented-call data understands callers the generic American-English bots fail on. Critical for our realtor / multicultural markets, and impossible to copy.
- **Compliance + embedded workflow** — once we answer the phones, write to their systems, and run their recalls, switching us out means losing their entire front desk.

The key reframe: our edge is not "the AI" (commoditizing) — it's the vertical data, integrations, and trust built on top of it.

How We Reach the Market

Sequencing — focus, not spread

Stage	Vertical	Why
Pilot #0	PPC Guru (our agency)	Speed-to-lead on our own ad leads — fastest, cleanest proof with real money.
Vertical 1	Physiotherapy (our clinics)	Own the customers; after-hours first; compliance handled by our platform partner.
Vertical 2	Construction / home services	Highest value, no compliance burden, our GC clients as partners.
Vertical 3	Realtor / mortgage	Owner-direct, our ~10,000-strong network, fast penetration.

The owner-direct insight

For owner-operated businesses, the buyer *is* the person missing the calls — no gatekeeper, no committee. The pitch is emotional and immediate: *"go to your kid's game, sleep through the night — your AI never misses a lead."* Shortest sales cycle in B2B.

Channels

Our own businesses → practice-management software partnerships → owner networks and associations → low-difficulty SEO (the open keywords) → AI-assisted outbound → referrals. We start where our distribution is unfair and expand from proof.

The Stress Test — We Tried to Kill It First

Before committing, we red-teamed the idea under every pressure. A plan that survives its own worst-case is worth backing.

Pressure	How it holds up
Is the problem really burning?	Sell recovered revenue, proven on the client's own line. Kill-signal if <2 missed calls/day.
Reliability on live calls (biggest risk)	Start after-hours/overflow (low stakes); guardrails; human handoff; QA on 100% of early calls.
Churn (esp. realtors)	Anchor on higher-value verticals; integrations create lock-in. Realtors = proof, not core.
Competition / commoditization	Win on vertical depth + data, not the AI. Generic bots die; vertical systems compound.
Platform / big-tech risk	Own the vertical, data, and relationships they won't replicate; stay nimble; be acquirable.
Legal / compliance	Inbound + consented speed-to-lead only — never cold outbound. Compliant by design.

Verdict: the idea survives — conditionally. The failure modes are real but *controllable by our own design choices* (after-hours-first de-risks reliability; higher-value verticals de-risk churn). That is exactly the kind of risk we want: the kind we control.

Architecture & Build Approach

We rent the commoditizing hard parts and build the defensible parts — and we architect from day one to swap providers as the technology evolves.

Phase 1 — ship fast

Build on **Retell** (turnkey voice orchestration with calendar/CRM integrations and healthcare compliance included for free), combined with best-in-class speech, language (Claude), and voice models, all behind a **provider-abstraction layer** so nothing locks us in.

Phase 2 — own the best tool

Migrate the core to open frameworks (**LiveKit / Pipecat**) to fully own interruption handling, a multi-vendor "best-of-breed" router, and our fine-tuned accent model — the last 10% that makes it the *best* product, not just a good one.

Design principle: the live call stays lightning-fast and lean; the heavy AI validation and learning happen *offline* after the call. The caller never waits, and the system gets smarter every night.

The lean, AI-leveraged team

One human (Shrikaanth) orchestrates ~8 internal **AI executor agents** (development, onboarding, quality-assurance, follow-up, monitoring) that run the repetitive work, with the human as reviewer and approver. We hire our first additional human only at ~30–40 customers. This is how a tiny team builds and runs a real product.

Compliance & Risk — Handled by Design

We mapped the full US + Canada regulatory landscape before building, so compliance is an architecture choice, never a retrofit.

Cross-cutting (handled once, centrally)

Calling/texting consent (TCPA / CASL), call-recording consent, caller-ID authentication, privacy law, **voiceprint/biometric consent** (a voice-specific landmine), and AI disclosure — built into the core so every vertical inherits it.

Per-vertical, sequenced by burden

Burden	Verticals	Our handling
Low	Home services, auto, restaurants	Cross-cutting only — our clean scaling lane.
Medium-High	Healthcare (physio), real estate	Platform provides HIPAA/BAA free; no-clinical-advice + emergency guardrails.
High (later)	Insurance, mortgage, finance	AI qualifies only — never advises/binds; anti-bias review built in.

The two golden rules: (1) the AI never gives regulated advice — it qualifies, schedules, informs, and hands off; (2) qualification logic is reviewed for anti-bias before any finance/real-estate launch.

Business Model & Unit Economics

70-80%

GROSS MARGIN

\$300-500

SETUP FEE (INSTANT CASH)

Appts booked

NORTH STAR METRIC

Pricing

A tiered subscription (\$199 / \$399 / \$799 per location per month) with included minutes and overage, a one-time setup fee, and — for the speed-to-lead use case — an optional **pay-per-booked-appointment** model (\$15–40) that aligns price directly to results and converts skeptics.

Margins

Delivery cost is ~\$0.10–0.18 per minute; at typical call volumes that's ~\$75–225/month against \$300–400 in revenue — **~70–80% gross margin**. The setup fee is near-pure margin. Per-appointment pricing runs ~85% margin and customers love it because they only pay for outcomes.

North Star

We measure one thing above all: **qualified appointments booked per month**. It is the value the customer pays for, it correlates with retention and expansion, and it keeps us honest — we only win when the customer wins.

The Roadmap — Phase 0 to Scale

Phase	What happens	Timing
Phase 0 — Foundation	Accounts, provider-abstraction layer, data-capture schema; sign vendor compliance agreements; draft consent/disclosure/data-rights.	Week 0–1
Phase 1 — Agency pilot	Speed-to-lead bot on PPC Guru's own ad leads; measure conversion lift.	Week 1–2
Phase 2 — Physio inbound	"Physio Brain", calendar integration, deploy to one clinic after-hours; QA every call.	Week 2–4
Phase 3 — Self-improving loop	Offline validation, executor agents, <1-day auto-onboarding.	Week 3–6
Phase 4 — Go live + first external customer	Full receptionist, billing, insurance; first paying external clinic.	Month 2
Phase 5 — Scale	Construction vertical, white-label, begin self-build migration.	Month 3+

Each phase has explicit **Build · Monitor · Compliance** gates (detailed in the companion Build Bible), sequenced so we never block first revenue and never retrofit compliance.

Investment, Status & Next Step

What it costs to get to a paying external customer

\$5-15k

TO FIRST PAYING CUSTOMER

4-6 wks

TO MVP

\$500-1.6k/mo

TO RUN ALL 7 OF OUR BUSINESSES

Infrastructure and tools run a few hundred dollars a month during build; legal setup (consent, terms, data-rights, compliance agreements) is a one-time \$1,500–5,000; the rest is our own time. We bootstrap on our businesses' cash to product-market fit, then raise from a position of proof.

Where we stand today

- **Research:** complete — 21-tab Data Bible with demand, winnability, and competitor evidence.
- **Strategy:** locked — product, vertical sequence, moat, pricing, and stress-test all decided and documented.
- **Plan:** a full Build Bible with architecture, tool selection, compliance register, and the phased roadmap.
- **Next step:** execute **Phase 0** and the agency pilot.

What we're asking of you, partner: alignment on this direction, and a green light to start Phase 0. The capital required is small; the evidence behind the decision is large; and the unfair advantages are already in our hands.

"We didn't pick this idea — the data did. Our job now is execution, and we've already de-risked the hardest parts."

Evidence & Source Files

Everything in this narrative is backed by committed, auditable artifacts in our repository:

Artifact	What it contains
SaaS Gap Research — Data Bible (21 tabs)	All demand (USA/Canada/World), seasonality, 4-yr trends, CPC, the 164-competitor teardown, complaint coding, and the Opportunity Score.
Demand keyword research + Ubersuggest set (560 kw)	Non-branded demand with volume, CPC, competition, and SEO difficulty.
Build Bible	Architecture, tool matrix, build journey (Build/Monitor/Compliance), the US+Canada regulatory register, and the AI agent org.
Strategy Findings & Recommendation	The decision brief and the grey-space analysis.

Method note: keyword volumes are bucketed by the source tools and should be read as ranges; market sizes are planning estimates; competitor specifics were cross-validated across two AI engines and the highest-confidence findings are those both engines independently agreed on. All figures are decision-grade, not audited financials.

Prepared June 2026 · Confidential · For partner discussion only.